

Elbasan Pulse

Everything the team needs for the 9-minute ideation checkpoint: the script, the scope mentors will test, and the answers to their questions.

15:45

CP1 slot opens, Friday 25 Sep
(Agenda)

9 min

per team, no slides

30 h

scope must be buildable in this
time

"Most dashboards assume the data is already unified. In a municipality, that assumption is the problem. Elbasan Pulse makes integration the product."

Team

FirmatGroup

Track

D: Municipal Data Integration & Decision-Support Dashboard

Challenge owner

Municipality of Elbasan

Sources

AI4Society 3-Day Agenda; Selected Participants' Guide; project design spec

1. At a glance

Format conflict: plan for 9 minutes. The Event Agenda puts CP1 at **15:45-17:00, 9 minutes per team, "Ideation", no slides**. The Participants' Guide says 18:30 and 1 minute, but calls its times provisional. Prepare the 9-minute talk; keep the 1-minute version (last page) ready in case the organisers confirm the shorter slot.

Problem

Elbasan's data is split across department systems, so leadership has no single view for planning and monitoring.

Users

Leadership (Mayor's cabinet), department heads, and planning and reporting staff.

Hypothesis

If onboarding a new data source takes minutes instead of an IT project, the municipality will keep one shared view up to date.

What mentors check at CP1

Pass condition

Mentors confirm the scope is **buildable in 30 hours**. Oversized scopes are cut on the spot.

So we finish with a short must-have list, a stretch list, and an explicit "not doing" list, and we name the cut order ourselves.

What we must show

- The problem, in the municipality's own words
- The user, with one concrete task each
- The solution hypothesis and the core flow
- An honest scope and data/AI boundaries

Our three claims

1. **Onboarding is minutes, not an IT project.** AI proposes column mappings; a person confirms; the mapping is reused next month.
2. **Every answer is auditable.** Each answer shows the SQL and the source file behind it.
3. **It transfers.** Any municipality, such as Gjakova, plugs in its own exports; the model and mappings are reusable.

Before you walk in

- ☐ Read the data package README (handed out 14:50-15:45) and fill every [yellow blank] in the script
- ☐ Assign a speaker to each block (section 3) and rehearse once with a timer
- ☐ Agree who notes mentor feedback and scope cuts during Q&A
- ☐ Decide: show the pre-event UI mockup only if asked, and say it is declared in Annex E

2. The 9-minute plan

Problem	Users	Solution	Core journey	AI	Scope	Data	Next	
0:00	1:00	2:00	3:30	5:00	6:15	7:30	8:15	9:00

Time	Block	One-line goal	Speaker	Scores on
0:00-1:00	Problem	The jury recognises Elbasan's own problem	_____	Relevance 20%
1:00-2:00	Users	Three users, one real task each	_____	User logic 15%
2:00-3:30	Solution	Integration is the product; state the hypothesis	_____	Innovation 10%
3:30-5:00	Core journey	Upload → confirm → see → ask → verify	_____	Feasibility 20%
5:00-6:15	AI	Where AI adds value, and its limits	_____	Use of AI 15%
6:15-7:30	Scope	Must-have, stretch, not doing	_____	Feasibility 20%
7:30-8:15	Data & honesty	Organiser data only; pre-event work disclosed	_____	Rules §10.2-10.3
8:15-9:00	Scale & next	Other municipalities; 4-week mentoring plan	_____	Scale 10% · Next 10%

Delivery tips

- About 130 words per minute: the script is about 1,150 words.
 - Open and close in Albanian; speak the body in English (the working language).
 - Hand-offs are one sentence: "Now [name] will walk you through the flow."
- If you are behind at 5:00, shorten the AI block to its two jobs plus "a person always confirms".
 - Never cut the scope block: it decides whether you pass.
 - End by inviting feedback on the scope; mentors like to be asked.

3. Script by speaker

Highlighted words are the ones to stress. Yellow blanks must be filled from the data package before the talk.

0:00

1 minute

Speaker: _____

Municipal relevance

1 • Problem

"Mirëdita. We are FirmatGroup, and we chose Track D. The Municipality of Elbasan told us something simple: its information sits in **separate department systems**, so nobody gets one view of services, operations and citizen requests. Planning and performance monitoring suffer. In practice, a department head exports a spreadsheet, someone in the Mayor's office copies numbers into another one, and by the time a report reaches leadership it is **late**, and no one can check **where a number came from**. We already see this in the data package: **[N files from N departments, each with different columns and formats]**."

Must land: the brief's words ("fragmented", "unified view", "planning and monitoring") plus one fact from the package.

1:00

1 minute

Speaker: _____

User logic

2 • Users

"We design for three users, exactly the ones in the brief. **Leadership**: the Mayor's cabinet wants to see in one glance how the municipality is doing, and to ask a question without waiting for an analyst. **Department heads**: they want their own performance against a target, and to see whether the backlog is growing. **Planning and reporting staff**: they do the integration by hand today. Their task is: 'a new export arrived, get it into the shared picture'."

2:00

1.5 minutes

Speaker: _____

Innovation

3 • Solution hypothesis

"Most dashboards assume the data is already clean and unified. In a municipality, **that assumption is the problem**. So Elbasan Pulse makes **integration the product**. Any department export, CSV or Excel, is onboarded in minutes: the system profiles the columns, AI proposes how each one maps to one shared data model, and **a person confirms**. The mapping is saved, so next month's export loads with one click. Once the data is unified, everything else becomes cheap: KPIs, department scorecards, trends, a monthly briefing, and an open API. Our hypothesis: if onboarding a new data source takes **minutes instead of an IT project**, the municipality will actually keep one shared view up to date."

3:30

1.5 minutes

Speaker: _____

User logic • Feasibility

4 • Core journey (the CP2 demo)

"Here is the one flow we will have working end to end at Checkpoint 2. **One**: a reporting clerk uploads **[real file name]**. **Two**: Pulse shows each column with a sample value and the AI suggestion, for example **['drejtoria' means department]**. The clerk corrects anything wrong and confirms. **Three**: the rows land in one database, and every row remembers which file it came from. **Four**: the dashboard updates: requests received, resolved on time, overdue, by department and month, against a target. **Five**: the department head asks, 'Which department has the most overdue requests?' Pulse answers and shows the **SQL and the source file**, so anyone can check it. That is the demo: **upload, confirm, see, ask, verify**."

5:00

1.25 minutes

Speaker: _____

Use of AI

5 • AI: value and limits

"AI does two jobs where it saves real time. **Column mapping**: a small, cheap model suggests mappings. It sees only column names, types and a few non-personal sample values, never personal data, and a person always confirms. **Questions in Albanian or English**: a stronger model turns the question into SQL, and that SQL is checked before it runs: read-only, one statement, approved tables only, limited rows. If the check fails, we say **'I can't answer that'** instead of guessing. What AI does **not** do: KPIs and trend alerts are plain calculations. If the AI budget runs out, the dashboard, alerts and API keep working. We track spend against our **[USD 25-50]** budget."

6:15

1.25 minutes

Speaker: _____

Feasibility (pass condition)

6 • Scope for 30 hours

"We want to be precise about the weekend. **Must-have by Checkpoint 2**: upload-and-map, the shared data model on the organisers' package, the KPI dashboard with filters, and the question box that shows its SQL. **Stretch by Checkpoint 3**: trend and anomaly alerts, an AI-written monthly briefing, and a map if the data has locations. **Not doing**: logins, live connections to municipal systems, writing back into them, and a mobile app. Those are next steps with the municipality. The technical base, a Python API with an embedded analytics database and a bilingual web app, is ready, so the 30 hours go into **integration logic, not setup**."

7:30

45 seconds

Speaker: _____

Rules §10.2, §10.3

7 • Data, privacy, honesty

"We use **only the organisers' data**. Personal fields are detected and never sent to the AI. Questions can only read, never change data. We also want to be transparent now, not on Sunday: before the event we prepared a generic project skeleton and a **visual mockup on clearly labelled fake numbers**, declared in our Annex E. Everything that makes it work, the ingest, the data model, the metrics and the copilot, gets built here on the real package."

8:15

45 seconds

Speaker: _____

Scalability • Next steps

8 • Scale and next steps

"Why this scales: the shared model and saved mappings are reusable. **Gjakova, or any municipality**, uploads its own exports and gets the same dashboard, and the API lets other systems, such as a citizen-reporting app, send data in. In the **four-week mentoring phase** we would test it with two Elbasan departments on their real monthly exports, measure onboarding time and data freshness, and prepare a pilot. Faleminderit. We'd love your feedback on the scope."

4. Scope card: buildable in 30 hours

This is what mentors will test. Put a name on every row before CP1.

Tier	Deadline	Deliverable	Estimate	Owner
Must-have	CP2 · Sat 15:00	Upload CSV/XLSX → profile columns → AI mapping proposal → confirm → load into DuckDB with source lineage	6 h	_____
		Canonical data model adapted to the data package	2 h	_____
		Metrics API; dashboard switched from sample data to real data	4 h	_____
		Ask box: question → guarded, read-only SQL → answer + SQL + source	5 h	_____
Stretch	CP3 · Sat 21:00	Signals: z-score / rolling-baseline alerts on metric series	3 h	_____
		AI-written monthly briefing (sq/en)	2 h	_____
		Map view (only if the data has locations)	2 h	_____
Not doing	n/a	User accounts / SSO · live connectors to municipal systems · write-back · mobile app		

Estimates are team-hours for planning, not promises; adjust them after reading the package.

If mentors say "too big": cut in this order

1. Map
2. AI briefing (keep the template briefing, which already works)
3. Signals (keep the fixed target rule, ≥ 90% on time)
4. XLSX support (CSV only)

What is already in place (disclosed)

- FastAPI + DuckDB skeleton with health check
- Claude client with budget cap and safe failure
- Bilingual web app shell; UI mockup on synthetic data
- CI, Docker, tests

No Track D logic: that is what we build here.

Team roles (Annex F categories)

Role	Weekend responsibility	Name
Technical	Ingest, warehouse, copilot, SQL guard	_____
Design / product	Mapping UX, dashboard wiring, demo flow	_____
Multidisciplinary (civic / domain)	Data package, KPI definitions, municipal Q&A	_____
Communication / business	Pitch, deck (max 8 slides), Annex E, next-steps plan	_____

5. How the talk maps to the jury scorecard

The final jury scores each criterion 1-5 (weighted total out of 100). Mentors at CP1 are already thinking in these terms.

Criterion	Weight	A top score looks like	Our evidence
Municipal relevance	20%	The challenge owner recognises their problem in the demo	Brief's own words and users; built on Elbasan's data package; Deputy Mayor attends the event
Feasibility	20%	Working prototype, credible assumptions, can continue	One small end-to-end flow; scaffold ready; AI failures never break the app
Use of AI	15%	AI adds genuine value; honest about limits	AI only where it saves hours; a person confirms; SQL shown; "can't answer" instead of guessing
User logic	15%	A real clerk could use it end to end	One flow per user: clerk uploads, head monitors, leadership asks
Innovation	10%	A better approach, not a re-skin	Integration-first and auditable answers, not another static dashboard
Scalability	10%	Transferable to other municipalities	Reusable model and mappings; open API; Gjakova as the example
Clarity of next steps	10%	Concrete plan for mentoring and piloting	Two departments, real monthly exports, measure onboarding time, pilot

Tie-breakers are municipal relevance first, then challenge-owner preference, then the technical jury's vote. Keep every sentence tied to Elbasan.

Rules that affect CP1

Rule	What it says	What we do
\$10.2	Core built during the event; pre-existing non-core parts allowed if disclosed; AI tools declared in submission and presentation	Mention the pre-event mockup and AI tools in the talk; keep Annex E current
\$10.3	Only organiser-provided or approved data; no personal data (risk of disqualification)	Swap synthetic data for the package first; ask a mentor to approve placeholders
\$10.5	Sun 11:00: deck (≤ 8 slides, template), demo link, tagged release, AI declaration	Late teams pitch last and lose next-steps points: tag early

6. Mentor questions and our answers

Likely question	Answer (keep it under 20 seconds)
What if the AI maps a column wrong?	A person confirms every mapping. Required fields are validated, and type errors are reported per column with row counts.
What if it generates bad or dangerous SQL?	A SQL guard allows only one read-only query over approved tables, with a row limit and a 5-second timeout. Otherwise it says "can't answer". The SQL is always shown.
How is this different from Power BI or Excel?	Those assume clean, modelled data and an analyst. Pulse does the integration step with AI and lets non-analysts ask questions in Albanian, with answers they can verify.
What did you build before the event?	A generic scaffold and a UI mockup on fake, labelled data, declared in Annex E. No Track D logic; the git history proves the timing.
Is 30 hours realistic?	The setup is done; the must-have list is four items; and we have a cut order ready.
What personal data does the AI see?	None. It sees column names, types and a few non-personal samples. Columns that look personal are excluded automatically.
What does it cost to run?	Open-source stack on one small server. AI costs a few cents per mapping or question, with a hard budget cap.
Why would the municipality keep using it?	Because monthly reporting gets faster: saved mappings make next month's upload one click, and the briefing writes itself from the data.
What happens after the hackathon?	Four weeks of mentoring: two departments, real monthly exports, measure onboarding time and freshness, then a pilot proposal.

Questions we ask

Mentors (at CP1) 1. Is our must-have list right for 30 hours, or should we cut further? 2. Which package files best show "fragmented systems" for the demo? 3. May we show the disclosed pre-event UI mockup?	Municipality (Sat 09:30 ICT talk, Q&A desk 10:00-13:00) 1. Which systems do departments export from today? 2. Which three numbers does leadership ask for most? 3. Weekly or monthly reporting, and who compiles it? 4. Are there official service targets (days to resolve)?
--	--

Feedback log (fill in during CP1)

Mentor / topic	Feedback	Action / owner

7. Data package worksheet

The package (one per track) holds anonymised samples, forms, scraped website content, category lists, the challenge statement and a README. Fill this in the first 20 minutes after hand-out; it feeds the script blanks and the canonical model.

File	Department / system	Rows	Key columns (→ canonical field)	Dates? Location?

Canonical fields to look for

Canonical table	Fields
request	id, category, department, location, created_at, closed_at, status, channel
department / service_category	name (sq/en), parent department
location	name, lat, lon, zone
budget_line	department, period, planned, actual
kpi_target	department, metric, period, target

Personal-data check (§10.3)

- ☐ No names, phones, emails or ID numbers passed to the AI
- ☐ Personal columns flagged and excluded from samples
- ☐ Only package files are used

Script blanks to fill

- ☐ Block 1: number of files and departments
- ☐ Block 4: real file name
- ☐ Block 4: a real column mapping example
- ☐ Block 5: our AI budget

8. Weekend timeline and gates

When	What	What we need ready
Fri 14:50-15:45	Track confirmation; data package, API keys, repo handed over	Read the README; fill the worksheet; test the API key
Fri 15:45-17:00	CP1 Ideation (9 min, no slides)	This pack; scope card with names
Fri 17:00-18:30	Team discussion and preparation	Apply mentor cuts; finalise the canonical model
Sat 09:00	Stand-up (30 seconds per team)	"Today: ingest and dashboard on real data by 15:00"
Sat 09:30	Municipal ICT presentation	Our municipal questions (section 6)
Sat 10:00-15:00	Build rounds 1-2; mentor rotation; municipal Q&A desk	Must-have items; use the help board if stuck for more than 1 hour
Sat 15:00-17:00	CP2 Core Demo (scheduled per team)	Upload → confirm → see → ask → verify, end to end
Sat 17:00	Pitch-craft session	Draft deck outline (8 slides)
Sat 21:00	CP3 Scope Freeze	Final feature list + who does what; no new features after this
Sun 09:00-11:00	Polish; mandatory dry-run with a mentor (10 min)	Rehearsed 5-minute pitch; demo on a stable URL or localhost
Sun 11:00	Submission deadline	Deck · demo link · tagged release · Annex E
Sun afternoon	Pitches: 5 min + 3 min jury Q&A	Results announced after both the Belgrade and Tirana hackathons

After the event (if we win)

October 2026: 4-week online mentoring (Reputeo) and the MVP Validation Lab in Tirana. February 2027: final presentation to the Municipality of Elbasan. Winning teams offer the municipality preferential terms (for example discounted pricing), agreed separately.

9. Pocket cue card

Print this page, cut along the dashed line, and hold it while speaking.

0:00	Problem: separate systems → no unified view → late reports, numbers nobody can trace. Package fact: 
1:00	Users: leadership (glance and ask) · dept heads (targets, backlog) · reporting staff (onboard exports)
2:00	Solution: "that assumption is the problem" → integration is the product → AI maps, a person confirms, mapping reused
3:30	Journey: upload → confirm → see → ask → verify (SQL + source)
5:00	AI: mapping (small model, no PII) · questions → guarded SQL · "can't answer" not guessing · no AI → still works
6:15	Scope: CP2 ingest, model, dashboard, ask · CP3 signals, briefing, map · NOT logins, live links, write-back, mobile
7:30	Honesty: organiser data only · pre-event scaffold and mockup declared in Annex E
8:15	Scale: Gjakova or any municipality · API · 4 weeks: 2 departments, measure, pilot · "Faleminderit, feedback on scope?"

1-minute fallback (if the Guide's format applies)

EN: "Elbasan's data sits in separate department systems, so leadership can't see the whole picture for planning and monitoring. Our users are the Mayor's office, department heads, and planning and reporting staff. Elbasan Pulse makes integration the product: upload any department export, AI proposes how it maps to one shared data model, and a person confirms it. On top sit KPIs, trend and anomaly signals, and a copilot that answers in Albanian or English, always showing its source and SQL. There's an open API so other systems, and other municipalities, can plug in."

SQ: "Në Bashkinë e Elbasanit, të dhënat janë të shpërndara nëpër sisteme të ndryshme të drejtorive. Kështu, drejtuesit nuk e shohin dot pamjen e plotë për planifikim dhe monitorim. Përdoruesit tanë janë zyra e Kryetarit, drejtuesit e drejtorive dhe stafi i planifikimit e raportimit. Elbasan Pulse e bën vetë integrimin e të dhënave produktin kryesor. Ngarkoni çdo eksport nga një drejtori, dhe AI-ja propozon si të lidhet me një model të përbashkët të dhënash. Një njeri e konfirmon. Mbi këtë ndërtojmë tregues kryesorë, sinjale për trendet dhe anomalitë, si dhe një asistent që u përgjigjet pyetjeve në shqip ose anglisht. Çdo përgjigje tregon burimin dhe pyetjen SQL pas saj. Kemi edhe një API të hapur, që sistemet e tjera, dhe bashkitë e tjera, të lidhen lehtësisht."